

Which Team will Win in League of Legends?

Name(s): Trevan Nguyen

Website Link: <https://sudasure.github.io/LeagueOfLegendsModel/>

Code

```
In [1]: import pandas as pd
import numpy as np
import os

import plotly.express as px
pd.options.plotting.backend = 'plotly'

from sklearn.model_selection import train_test_split
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.ensemble import RandomForestClassifier
from sklearn.preprocessing import StandardScaler, QuantileTransformer, OneHotEncoder
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import precision_score
from sklearn.utils import resample
from sklearn.metrics import classification_report
```

Framing the Problem

Prediction Problem: Predict if a team will win or lose a game based on post-game data.

Type: Binary Classification

Response Variable: Game result (Win/Loss)

Justification: Predicting the outcome of a League of Legends game can be useful for analyzing team performance and making strategic decisions. By predicting whether a team will win or lose based on post-game data, we can gain insights into factors that contribute to a team's success.

Evaluation Metric: Accuracy

Justification: Accuracy is appropriate for this classification problem as it measures the overall correctness of the predictions. We want to accurately classify whether a team will win or lose a game to evaluate the model's predictive performance.

```
In [2]: # Load League dataframe
lol_raw = pd.read_csv('2022_LoL_esports_match_data_from_OraclesElixir.csv', low_memory=False)
lol_raw.head()
```

Out[2]:

	gameid	datacompleteness	url	league	year	split	playoffs	date	game	patch	...	opp_csat15	golddiffat15	xpdiffat15	csdiffat15	killsat15	assistsat15	deathsat15	opp_killsat15	opp_assistsat15	opp_deathsat15
0	ESPORTSTMNT01_2690210	complete	NaN	LCK CL	2022	Spring	0	2022-01-10 07:44:08	1	12.01	...	121.0	391.0	345.0	14.0	0.0	1.0	0.0	0.0	1.0	0.0
1	ESPORTSTMNT01_2690210	complete	NaN	LCK CL	2022	Spring	0	2022-01-10 07:44:08	1	12.01	...	100.0	541.0	-275.0	-11.0	2.0	3.0	2.0	0.0	5.0	1.0
2	ESPORTSTMNT01_2690210	complete	NaN	LCK CL	2022	Spring	0	2022-01-10 07:44:08	1	12.01	...	119.0	-475.0	153.0	1.0	0.0	3.0	0.0	3.0	3.0	2.0
3	ESPORTSTMNT01_2690210	complete	NaN	LCK CL	2022	Spring	0	2022-01-10 07:44:08	1	12.01	...	149.0	-793.0	-1343.0	-34.0	2.0	1.0	2.0	3.0	3.0	0.0
4	ESPORTSTMNT01_2690210	complete	NaN	LCK CL	2022	Spring	0	2022-01-10 07:44:08	1	12.01	...	21.0	443.0	-497.0	7.0	1.0	2.0	2.0	0.0	6.0	2.0

5 rows × 123 columns

```
In [3]: # Get the columns that will be useful for our analysis
lol = lol_raw[['gameid', 'datacompleteness', 'side', 'position', 'result', \
              'kills', 'deaths', 'assists', 'pentakills', 'firstblood', 'team kpm', \
              'dragons', 'firstherald', 'firstbaron', 'firsttower', 'towers', 'damagetochampions', 'visionscore', 'totalgold']]

# Extract the team data and fill NaN values with 0.0 because they are missing because the value would have been 0.0 anyways
lol_teams = lol[lol['position']=='team'].fillna(0.0)
lol_teams.head()
```

Out[3]:

	gameid	datacompleteness	side	position	result	kills	deaths	assists	pentakills	firstblood	team kpm	dragons	firstherald	firstbaron	firsttower	towers	damagetochampions	visionscore	totalgold
10	ESPORTSTMNT01_2690210	complete	Blue	team	0	9	19	19	0.0	1.0	0.3152	1.0	1.0	0.0	1.0	3.0	56560.0	197.0	47070
11	ESPORTSTMNT01_2690210	complete	Red	team	1	19	9	62	0.0	0.0	0.6655	3.0	0.0	0.0	0.0	6.0	79912.0	205.0	52617
22	ESPORTSTMNT01_2690219	complete	Blue	team	0	3	16	7	0.0	0.0	0.0851	1.0	1.0	0.0	0.0	3.0	59579.0	277.0	57629
23	ESPORTSTMNT01_2690219	complete	Red	team	1	16	3	39	0.0	1.0	0.4541	4.0	0.0	1.0	1.0	11.0	74855.0	346.0	71004
34	8401-8401_game_1	partial	Blue	team	1	13	6	35	0.0	0.0	0.5714	2.0	0.0	0.0	0.0	8.0	40086.0	162.0	45468

Baseline Model

Features: First blood and first tower

Used RandomForestClassifier and OneHotEncoder for these nominal categorical variables The ‘firstblood’ indicates whether the team achieved the first kill in the game and ‘firsttower’ indicates whether the team destroyed the first tower. The performance of the model is evaluated using accuracy.

```
In [4]: # Select features and target variable
features = lol_teams[['firstblood', 'firsttower']]
target = lol_teams['result']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(features, target, test_size=0.2, random_state=42)

# Define the preprocessing steps for the categorical columns
preprocessor = ColumnTransformer(
    transformers=[
        ('cat', OneHotEncoder(), ['firstblood', 'firsttower'])
    ])

# Create the baseline model pipeline
pipeline = Pipeline([
    ('preprocessor', preprocessor),
    ('classifier', RandomForestClassifier())
])

# Train the model
pipeline.fit(X_train, y_train)

# Evaluate the model on the test set
accuracy = pipeline.score(X_test, y_test)
y_pred = pipeline.predict(X_test)
accuracy = pipeline.score(X_test, y_test)
report = classification_report(y_test, y_pred)

print("Classification Report:")
print(report)
```

Classification Report:					
	precision	recall	f1-score	support	
	0	0.63	0.74	0.68	2465
	1	0.70	0.58	0.63	2515
accuracy				0.66	4980
macro avg	0.66	0.66	0.66		4980
weighted avg	0.67	0.66	0.66		4980

Final Model

Features: First blood, first tower, damage to champions, and kills

Used RandomForestClassifier, OneHotEncoder, StandardScaler, QuantileTransformer, GridSearchCV, etc. The added features ‘damagetochampions’ indicates the amount of total damage done to enemy champions and ‘kills’ how many takedowns a given team achieved. The performance of the model is evaluated using accuracy.

Best hyperparamaters are the second/middle options

```
In [5]: # Select the features and the target variable
X = lol_teams[['firstblood', 'firsttower', 'damagetochampions', 'kills']] # Features
y = lol_teams['result'] # Target variable

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Define the feature transformation steps
preprocessor = ColumnTransformer(
    transformers=[
        ('cat', OneHotEncoder(), [0]) # Categorical feature to be one-hot encoded
    ],
    remainder='passthrough' # Pass through the remaining numerical features as-is
)

# Define the final model pipeline
pipeline = Pipeline([
    ('preprocessor', preprocessor),
    ('feature_engineering', ColumnTransformer(
        transformers=[
            ('num1', StandardScaler(), [2]), # StandardScaler on feature3
            ('num2', QuantileTransformer(), [3]) # QuantileTransformer on feature4
        ],
        remainder='passthrough'
    )),
    ('classifier', RandomForestClassifier()) # Random Forest classifier as an example
])
```

```
# Define the hyperparameters to tune
param_grid = {
    'classifier__n_estimators': [10, 20, 30], # Number of trees in the forest
    'classifier__max_depth': [None, 5, 10], # Maximum depth of the tree
}

# Perform grid search with cross-validation to find the best hyperparameters
grid_search = GridSearchCV(pipeline, param_grid, cv=5)
grid_search.fit(X_train, y_train)

# Get the best hyperparameters and model
best_params = grid_search.best_params_
final_model = grid_search.best_estimator_

# Evaluate the model on the test set
y_pred = final_model.predict(X_test)
accuracy = final_model.score(X_test, y_test)
report = classification_report(y_test, y_pred)

print("Classification Report:")
print(report)
print("Best Hyperparameters:", best_params)
```

Classification Report:					
	precision	recall	f1-score	support	
	0	0.89	0.79	0.84	2465
	1	0.82	0.90	0.86	2515
	accuracy			0.85	4980
	macro avg	0.85	0.85	0.85	4980
	weighted avg	0.85	0.85	0.85	4980

Best Hyperparameters: {'classifier__max_depth': 5, 'classifier__n_estimators': 10}

Fairness Analysis

Null Hypothesis: The model is fair. The precision for the blue team and red team is roughly the same, and any differences are due to random chance.

Alternative Hypothesis: The model is unfair. The precision for the blue team is lower than the precision for the red team.

```
In [6]: # Assume X_test and y_test are the test features and labels
X_blue_team = lol_teams[lol_teams['side'] == 'Blue']
y_blue_team = lol_teams[lol_teams['side'] == 'Blue']

X_red_team = lol_teams[lol_teams['side'] == 'Red']
y_red_team = lol_teams[lol_teams['side'] == 'Red']

# Make predictions for blue team and red team separately
y_pred_blue_team = final_model.predict(X_test[y_test == 1])
y_pred_red_team = final_model.predict(X_test[y_test == 0])

# Calculate the precision score for each group
precision_blue_team = precision_score(y_test[y_test == 1], y_pred_blue_team)
precision_red_team = precision_score(y_test[y_test == 0], y_pred_red_team)

# Define the number of permutations
num_permutations = 100

# Initialize an array to store the permutation test statistics
perm_scores = np.zeros(num_permutations)

# Perform the permutation test
for i in range(num_permutations):
    # Permute the labels within each group
    perm_y_blue_team = resample(y_test[y_test == 1])
    perm_y_red_team = resample(y_test[y_test == 0])

    # Concatenate the permuted labels with the original labels
    perm_y = np.concatenate((perm_y_blue_team, perm_y_red_team))

    # Make predictions on permuted labels
    perm_pred_blue_team = final_model.predict(X_test[y_test == 1])
    perm_pred_red_team = final_model.predict(X_test[y_test == 0])

    # Calculate the precision score for permuted groups
    perm_precision_blue_team = precision_score(perm_y_blue_team, perm_pred_blue_team)
    perm_precision_red_team = precision_score(perm_y_red_team, perm_pred_red_team)

    # Calculate the test statistic (difference in precision)
    perm_scores[i] = perm_precision_blue_team - perm_precision_red_team

# Calculate the observed test statistic (difference in precision)
observed_diff = precision_blue_team - precision_red_team
```

```
# Calculate the p-value
p_value = np.mean(perm_scores >= observed_diff)

print("Precision for blue team:", precision_blue_team)
print("Precision for red team:", precision_red_team)
print("Observed difference:", observed_diff)
print("p-value:", p_value)
```

```
Precision for blue team: 1.0
Precision for red team: 0.0
Observed difference: 1.0
p-value: 1.0
```

Looks like we fail to reject the null in this case. The model seems to be fair. The precision for the blue team and red team is roughly the same, and any differences are due to random chance.